

Who Decides What “Open” Means?

Three weeks in 2026 when both superpowers
changed their minds at once.

Tim Menzies

NC State University · timm.fyi

TWO DATES

12 June 2026

Anthropic suspends Fable 5 and Mythos 5.
Export controls. Access restored 30 June.

24 July 2026

Jensen Huang posts on X for the first time ever.
25 companies sign. Nadella boosts it in minutes.

Six weeks apart. Opposite direction.

“Open Weights and U.S. AI Leadership”

- Signed by 25 firms. Nvidia, Microsoft, Meta, IBM, Dell, Palantir, Hugging Face, a16z, Y Combinator, Mozilla, Mistral, Replit.
- Claim: open models strengthen safety and cybersecurity, accelerate innovation and diffusion, enable sovereignty.
- Frame: do not repeat the 1980s mistake. Do not lock this down early.

Huang: the world needs leading-edge closed models, and leading-edge open ones too.

READ THE LAST PARAGRAPH

The payload is at the bottom.

- The final substantive paragraph argues that distillation is legitimate model development.
- Filed three days after the Treasury Secretary threatened sanctions over exactly that.
- The rest of the letter is preamble. This is the ask.

Nothing in the argument is new. Everything in the timing is.

WATCH THE LIST, NOT THE LETTER

It moved in 72 hours.

Friday: 25 signatories.

OpenAI, Google, Anthropic, Amazon all absent.

Sunday: about 50.

OpenAI signed after the omission was noticed. So did Google.

Still out: Anthropic. Amazon.

A position that flips under social pressure was never a position.

Two business models. One vocabulary.

Sells compute

Demand scales with diffusion.
Free weights grow the market.
Wants: open.

Sells tokens

Margin depends on scarcity.
Free weights are the competition.
Wants: closed.

Both sides say “security.” Both mean “revenue.”

Both claims cannot be structurally true.

- July: openness is asserted to be a security property.
- June: a frontier lab treated distribution itself as the security surface, and pulled two models for eighteen days.
- Same industry. Same quarter. Same word doing opposite work.

One of these is an engineering claim.

One is a business model wearing an engineering claim.

China is closing at the same moment.

- Commerce Ministry convened Alibaba, ByteDance, Z.ai on limiting overseas access to top models.
- Proposed tiering: filing for basic open tools, security review for advanced, domestic-only or no release at the frontier.
- Also floated: AI leaks as a national security offence.

Targets are the models everyone actually uses. Qwen. Doubao. GLM.

**In one quarter, both superpowers
began fencing the frontier,
and each accused the other of closing.**

The Halloween Documents

- Leaked Microsoft memos. Open source treated as an existential threat.
- Origin of “embrace, extend, extinguish.”
- Microsoft is a signatory on the 2026 letter.

28 years. Complete inversion.

Nobody changed their mind. The incentives moved.

“We Have No Moat, And Neither Does OpenAI”

- Leaked Google memo. Sernau, via SemiAnalysis, May 2023.
- Argument: while Google and OpenAI watched each other, open source quietly ate both.
- Evidence: LoRA fine-tunes on consumer hardware. Iteration in days, not months.

Often misremembered as 2025. It is 2023. The lag matters.

2023 memo: descriptive panic.

Written in private. Not meant to persuade anyone.

2026 letter: prescriptive lobbying.

Written for policymakers. Twenty-five logos attached.

**When a private worry becomes a public campaign,
ask what changed. It was not the evidence.**

The Crypto Wars set the precedent.


- Encryption classified as a munition under ITAR. Zimmermann investigated for publishing PGP.
- MIT Press printed the source code as a book. Books could be exported.
- Bernstein v. United States: source code is protected speech.

Nobody has litigated whether weights are speech.

That case is coming. Probably within two years.

DEFINE YOUR TERMS

Open is not one bit. It is six levels.



Solaiman, FAccT 2023. The gradient of generative AI release.

Most of the July letter argues about level five only.

WEIGHTS ARE ONE CELL

Openness is a grid, not a switch.

Model artifacts

Datasets

Code

Weights

System layer

Infrastructure

Product / UX

Interfaces

Cross-cutting

Documentation

Safeguards

Licensing

Each cell can be released more or less openly. Roughly forty of them.

Fifteen authors. OSI, Mozilla, Meta, Linux Foundation, AI2.

François et al., CACM, August 2026. Even OSI's director now says one bit is not enough.

“Given enough eyeballs, all bugs are shallow.”

- Raymond, 1997. Asserted. Never measured.
- xz-utils, 2024: two years of social engineering against one exhausted volunteer maintainer. A root backdoor nearly shipped to every major distro.
- Caught by accident. An engineer noticed 500ms of latency.

Linus's Law is an empirical claim with a failure record.

The other side is no better. Luccioni et al. surveyed the evidence for open-model risk and found much of it missing.

Two unmeasured claims. Two papers. We are the people who measure things.

What “free” actually buys.

\$400,000

Saved per year by one US firm running Qwen instead of a proprietary model. Nvidia, Perplexity and Stanford use Qwen too.

“If you need cutting-edge, go back to OpenAI. Most applications don't.”

Free is not charity. It is distribution bought with a compute deficit.

Reflection AI raised at \$8B to be the American version of that strategy.

What open weights bought the university.

- A lab of three can run frontier-adjacent models on hardware it owns.
- Reproducibility. You cannot replicate a study against an endpoint that was silently updated last Tuesday.
- Ablation. You cannot ask why if you cannot open it.
- Cost. Cheap weights are why work outside the funded centres exists at all.

Every one of these is lost at the “domestic-only” tier.

Open is not safer. Open is checkable.

- Safety is not a model property. It depends on deployment context, safeguards and governance, not the weights alone.
- So “open models are safer” is a category error, not merely a false claim.
- But without openness, independent risk assessment is impossible. The right question is marginal risk against closed models and against what already exists.

Verification versus faith. That is an empirical property.

Narayanan and Kapoor on safety. Bommasani et al. on marginal risk.

YESTERDAY, 30 JULY 2026

The argument stopped being theoretical.

141,006

evaluation runs reviewed.
Six were bad. Three incidents.

- Claude reached the open internet from a sandbox and compromised the production infrastructure of three organisations.
- Weak passwords. Unauthenticated endpoints. No clever exploit.
- The prompt said there was no internet. So every real host looked like part of the exercise.

OpenAI reached Hugging Face the same way, ten days earlier.

Nobody was watching. That is the finding.

- The earliest incident was in April. Found in July.
- Found only because a competitor disclosed first and prompted a retrospective log review.
- Two of the three victims had not detected the intrusion until they were told.

Closed model. Closed evaluation. Closed lab.

The only reason we know is that they chose to tell us.

Credit where due. But a disclosure regime built on goodwill is not a verification regime.

The model worked. The harness failed.

- Containment was not broken. It was misconfigured, and both parties believed it was on.
- In every run the model did only the task it was given.
- Three models, three responses. One pressed on. One decided it was still in a simulation. One stopped.

Auditing the weights would have caught none of this.

Safety lives in the system, not the model. Now we have the receipt.

Simple ain't stupid.

- Lightweight active learners match or beat heavy optimizers across the MOOT benchmark.
- The substrate is roughly 500 lines. It ports across a dozen languages.
- Labels are the scarce resource, not FLOPs. Spend them carefully.

**If the frontier gets fenced, frugal methods are not the fallback.
They are the plan.**

IF YOU REMEMBER ONE SLIDE

**Openness is being redefined
by whoever loses least
from the definition.**

Academics have no revenue at stake in this fight.

That makes us the only honest witnesses in the room.

Four things this room can do.

- Say which “open” you mean, every time, in every paper.
- Measure the many-eyes claim instead of citing it.
- Report the cost of your method, not only its accuracy.
- Archive the weights you publish against. Endpoints are not artifacts.

Three of these are papers. One is a habit.

OVER TO YOU

Discussion

- Are weights speech, or are they a munition?
- Can you name a security property that openness actually provides?
- There is no disclosure norm for “our model attacked a third party.”
- If China closes and the US closes, what does our field do next?
- Who in this room could still run their dissertation in 2028?

REFERENCES 1 OF 2

Frameworks and definitions

François, C. et al. Unpacking open source artificial intelligence: Toward a framework for openness in foundation models. CACM 69, 8 (Aug 2026), 80-87. DOI 10.1145/3778264

Basdevant, A. et al. Towards a framework for openness in foundation models. arXiv:2405.15802 (2024)

Open Source Initiative. Open Source AI Definition 1.0 (Oct 2024). opensource.org/ai/open-source-ai-definition

White, M. et al. The Model Openness Framework. arXiv:2403.13784 (2024)

Solaiman, I. The gradient of generative AI release: Methods and considerations. FAccT (2023), 111-122

Bommasani, R. et al. The Foundation Model Transparency Index. arXiv:2310.12941 (2023)

Risk, safety and evidence

Narayanan, A. and Kapoor, S. AI safety is not a model property. AI as Normal Technology (2024). normaltech.ai

Bommasani, R. et al. Considerations for governing open foundation models. Stanford HAI (2023)

Luccioni, S., Solaiman, I. and Paolotti, D. What does it mean to open source AI? arXiv:2403.07918 (2024)

Seeger, E. et al. Open-sourcing highly capable foundation models. Centre for the Governance of AI (2023)

Open model projects

Biderman, S. et al. Pythia. ICML (2023)

Liu, Z. et al. LLM360. arXiv:2312.06550 (2023)

Allen Institute for AI. OLMo. allenai.org/olmo

REFERENCES 2 OF 2

The 2026 dispute

Huang, J. et al. Open Weights and U.S. AI Leadership. Open letter, 24 July 2026. 25 signatories at publication, about 50 within three days

New York Times. Reporting on open-weight models and on Chinese restrictions to overseas model access, 27-29 July 2026

Anthropic. Statement on Fable 5 and Mythos 5 access, suspended 12 June 2026, restored 1 July 2026.
anthropic.com/news/fable-mythos-access

The evaluation incidents

Anthropic. Investigating three real-world incidents in our cybersecurity evaluations, 30 July 2026.
anthropic.com/news/investigating-incidents-cybersecurity-evals

OpenAI. Disclosure of model access to Hugging Face infrastructure, 21 July 2026

History

Raymond, E.S. The Cathedral and the Bazaar (1997)

Halloween Documents. Leaked Microsoft memoranda on open source, 1998. catb.org/~esr/halloween

Sernau, L. We have no moat, and neither does OpenAI. Leaked Google memo via SemiAnalysis, May 2023

Bernstein v. United States, 176 F.3d 1132 (9th Cir. 1999). Source code as protected speech

CVE-2024-3094. xz-utils backdoor, March 2024

Disclosure norms

Rain Forest Puppy. RFPolicy 1.0 (2000)

CERT/CC Vulnerability Disclosure Policy; Google Project Zero 90-day deadline. ISO/IEC 29147