Response to Reviewers (Round 2)

Can AI be Easy? Lessons Learned from the EZR.py Toolkit

Menzies, Srinivasan, Ganguly — Software: Practice and Experience

We thank the reviewer for a second careful reading. All new or altered text in this revision is again printed in blue. Point by point:

Reviewer 1

M1(a). *The abstract's text-mining claim is contradicted by §7.4.* Fixed. The abstract now states that the recall match holds on three of the four corpora and that on the fourth (Kitchenham) the filter converges faster but plateaus below FASTREAD's ceiling, pointing to the fuller treatment in §7.

M1(b). *"Tested on the 120+ tasks" overstates the evidence for results (a) and (c).* Fixed. The abstract now scopes each result: (a) 20 randomly selected MOOT benchmarks, (b) the MOOT tasks, (c) Yu et al.'s four SLR corpora. The matching overstatement in §8.1 is likewise rescoped: the 124-task claim now attaches only to the active-learning results, and the optimizer and text-mining studies are scoped explicitly.

M1(c). *Tables 2 and 4 still present companion-derived numbers as findings of this paper.* Fixed. The §6 row of Table 2 now opens "A companion study…(not a result of this paper) reports…", and the label-hunger row of Table 4 marks its comparator numbers as drawn from companion studies.

M2. *"Co-authored by the first author" understates the overlap now that Ganguly is an author.* Fixed. All instances (abstract, §6.4.4, §8.2, §8.3, §8.4, and the closing remarks) now read "co-authored by authors of this paper."

M3(a). *Cite the version-specific DOI and git tag.* Fixed, with an apology: DOI 10.5281/zenodo.17118009 was in fact already version-specific (it resolves to the v0.5.0 deposit); our text mislabeled it a concept DOI. The paper now names the tag and version-specific DOI at every point of use (v0.9.4, 10.5281/zenodo.22554751, freshly archived for this revision; listings trace to v0.5.0, 10.5281/zenodo.17118009) and gives the true concept DOI (10.5281/zenodo.11183058) for the full release history.

M3(b). *Confirm the Zenodo deposit contains the 400-line single-file version.* Confirmed. We downloaded the v0.5.0 deposit from Zenodo (timm/ezr-v0.5.0.zip under DOI 10.5281/zenodo.17118009) and verified it contains ezr/ezr.py: 412 physical lines, 44 functions, matching the listings of this paper. The current head of the repository (v0.9.4, two files, less than 1000 lines) is now also archived, under DOI 10.5281/zenodo.22554751, and the paper's size claims now say "less than 1000 lines" so that the text matches the current release. Additionally: a new Appendix A gives tested two-minute instructions (pip install ezr plus the MOOT data) for running the toolkit's fast demos and its full self-test suite, with sample output, so that an editor or reviewer can check the code runs without setting up our full experiments.

M4. *Table 4's "weeks" contradicts §2.6's "one session," undermining the order-of-magnitude claim.* Fixed. §2.6 now says "weeks of sessions vs. a semester course," and the order-of-

magnitude claim explicitly excludes the pedagogy axis, noting that the gap there is weeks vs. a semester and rests on author experience rather than formal evaluation.

M5(a). *§8.4's surrogate-bias paragraph is stale.* Fixed. The paragraph now names panel (a), corrects the count (sa-r at 100 on 13 of 20 rows, mean 98, under the new 100-repeat runs), reports the panel (b) view (no cell at 100, mean 94), and downgrades "decisive evidence" to "suggestive": the surrogate drives the search to within a few points of the optimum, but panel (b) shows the optimum is approached, not exactly attained.

M5(b). *The regenerated panel (a) shifts several column means; 20 repeats do not stabilise ls-r.* The caption of Table 5 now states that these runs supersede the original submission's table and that high-variance columns shift several points between regenerations. Further, we took the first of the reviewer's two offered remedies: the whole table (all four treatments, both panels, one set of runs) was regenerated with the repeat count raised from 20 to 100, so each cell mean now carries a standard error of sd/10. Under the 100-repeat runs the qualitative conclusions are unchanged (ls-r remains worst and most erratic; ls+r, mean 97, still does not pass SA, mean 98; restart on SA remains a rounding error), and all numbers quoted in §5.2 and §8.4 have been updated to the new runs.

Minor 1. *"?" bibliography entries; SHAP and ReliefF uncited.* Fixed. Five missing database entries (Lundberg & Lee 2017 for SHAP; Robnik-Šikonja & Kononenko 1997 for ReliefF; Yu et al. 2018 "Finding better active learners"; PEP 544; PEP 557) have been added, repairing all eight "?" instances.

Minor 2. *Broken cross-reference and unbalanced parenthesis in §6.2.* Fixed; the reference now resolves to Figure 1 and the parenthesis is balanced.

Minor 3. *The §7 numbers do not reconcile (warm start 24, budget 50, x-axes to 300–600, plateaus "after 100 labels," 40 labels/108 papers, "100 vs. 300+").* Fixed; the reviewer caught real errors, and we thank them. The true protocol (now stated in §7.3, and readable from the "40" and "Baseline" marks on Figure 5 itself) is: warm start of 40 labeled samples (20 known-relevant plus 20 presumptive negatives), not 24; then one label per acquisition, with each run continuing to the number of labels Yu et al. spent on that corpus (290–630), which is what the x-axis counts. The "budget was 50" sentence was a copy-slip from §6 and has been deleted. The Kitchenham claim is now stated in the same currency (after the 40-label warm start, recall is 87%, settling to 82% = 108 of 132). Yu et al.'s labeling cost is now given as 300–630 in both §7.2 and the §7.4 summary.

Minor 4. *Confirmed as fixed.* We thank the reviewer; no further action was needed on this point.

Minor 5. *Three identical affiliations.* Fixed; all three authors now share one affiliation.