

Retrospective: Data Mining Static Code Attributes to Learn Defect Predictors

Tim Menzies

September 23, 2025

The Portland Context

- Born from open source culture in Portland, Oregon
- *"We wore no suite and tie in our photos. We did not comb our hair"*
- Philosophy: `svn commit -m "share stuff"` will change SE research
- **Key Insight:** Walking around Chicago's Grant Park (2004)
 - **Tim Menzies** and **Jelber Sayyad** lamented: *"Must do better... Why don't we make conclusions reproducible?"*

The Radical Idea

- In 2025 hard to believe "reproducible SE" was radical
- **Lionel Briand** (2006): *"no one will give you data"*
- Yet we persisted...

Two-Part Vision:

- 1 **Annual conference** on predictor models in SE (to share results)
- 2 **Repository** of 100s of SE datasets: defect prediction, effort estimation, Github issue close time, bad smell detection

Growth Trajectory:

- Repository grew large; moved to **Large Hadron Collider** (Seacraft data at Zenodo)
- Research students ran weekly sprints scouring SE conference tables of content
- **Gary Boetticher, Elaine Weyuker, Thomas Ostrand, Guenther Ruhe** joined steering committee → prestige for growth

PROMISE vs MSR:

- **MSR**: Gathering initial datasets (**Devanbu [Dev15]**)
- **PROMISE**: Post-collection analysis, consistent data upload and re-examination [**Rob10**]

Early Results:

- Other areas struggled with reproducibility, we swam in an ocean of reproducibility
- Papers applied elaborate tool sets to COC81, JM1, XALAN, DESHARNIS datasets
- First decade: Numerous successful papers using consistent data re-examination

Research Question: Can data mining algorithms learn software defect predictors from static code attributes?

Why This Matters:

- *"Software quality assurance budgets are finite while assessment effectiveness increases exponentially with effort"* [Fu16]
- *"Software bugs are not evenly distributed across a project"* [Ham09], [Ost04], [Mis11]
- Defect predictors suggest where to focus expensive methods

Counter-Arguments Addressed:

- 1 *"Specific metrics matter"* (1990s heated debates: McCabe vs Halstead)
- 2 *"Static code attributes do not matter"* (Fenton & Pfleeger, Shepperd & Ince)

"Specific metrics do not always matter in all data sets. Rather, different projects have different best metrics."

Supporting Evidence:

- Feature pruning experiment on **3 dozen metrics across 7 datasets**
- Results: Pruning selected just **2-3 attributes per dataset**
- **No single attribute** selected by majority of datasets
- Different projects preferred different metrics (McCabe vs Halstead vs lines of code)
- Theoretical debates of 1990s (metric X vs metric Y) proven empirically unfounded

Menzies's Corollary:

"To mine SE data, gather all that can be collected (cheaply) then apply data pruning to discard irrelevancies."

Practical Impact:

- Changed SE data mining methodology from "careful metric selection" to "gather everything, prune later"
- Influenced thousands of subsequent studies using this approach

"Static code attributes do matter. Individually, they may be weak indicators. But when combined, they can lead to strong signals that outperform the state-of-the-art."

Supporting Evidence:

- **Fenton & Pfleeger:** Same functionality, different constructs → different measurements
- **Shepperd & Ince:** Static measures often "no more than proxy for lines of code"
- **Our Response:** Stress-tested these views by documenting baselines, then showing detectors from static attributes **much better** than baselines
- **Key Finding:** Multi-attribute models outperformed single-attribute models

Key Quote: *"Paradoxically, this paper will be a success if it is quickly superseded."*

Citation Impact:

- **2016:** Most cited paper (per month) in software engineering
- **2018:** 20% of Google Scholar Software Metrics IEEE TSE papers used PROMISE datasets [Men07]
- **Current:** 1924 citations (paper) + 1242 citations (repository)

Industrial Adoption:

- **Wan et al. [Wan20]:** 90%+ of 395 commercial practitioners willing to adopt defect prediction
- **Misirli et al. [Mis11]:** 87% defect prediction accuracy, 72% reduced inspection effort, 44% fewer post-release defects
- **Kim et al. [Kim15]:** Samsung Electronics API development - 0.68 F1 scores, reduced test case resources

Rahman et al. [Rah14] Comparison:

- **Static analysis tools:** FindBugs, Jlint, PMD
- **Statistical defect prediction:** Logistic regression models
- **Result:** *"No significant differences in cost-effectiveness were observed"*

Critical Advantage:

- Defect prediction: Quick adaptation to new languages via lightweight parsers
- Static analyzers: Extensive modification required for new languages
- **Implication:** Broader applicability across programming ecosystems

Extended Applications:

- **Security vulnerabilities** [Shi13]
- **Resource allocation** for defect location [Bir21]
- **Proactive defect fixing** [Kam16], [LeG12], [Arc11]
- **Change-level/just-in-time prediction** [Yan19], [Kam13], [Nay18], [Ros15]
- **Transfer learning** across projects [Kri19], [Nam18]
- **Hyperparameter optimization** [Agr18], [Che18], [Fu17], [Tan16]

Research Evolution:

- From binary classification to multi-objective optimization
- From release-level to line-level prediction (**Pornprasit et al. [Por23]** - TSE Best Paper 2023)

Phase Evolution:

- 1 *"Data? Good luck with that!"* - Resistance and skepticism
- 2 *"Okay, maybe it's not completely useless."* - Grudging acknowledgment
- 3 *"This is the gold standard now."* - Required baseline, field norms
- 4 *"A graveyard of progress."* - Stifling creativity, outdated paradigms

The Problem:

- Decade 2: Continued use of COC81 (1981), DESHARNIS (1988), JM1 (2004), XALAN (2010)
- **Editorial Policy Change:** Automated Software Engineering journal now desk-rejects papers based on 2005 datasets

Menzies's 3rd Law:

"Turkish toasters can predict for errors in deep space satellites."

Supporting Evidence: - **Transfer learning research [Tur09]:** Models from **Turkish white goods** successfully predicted errors in **NASA systems** - Expected: Complex multi-dimensional transforms mapping attributes across domains - **Reality:** Simple nearest neighboring between test and training data worked perfectly - **Implication:** *"Many distinctions made about software are spurious and need to be revisited"*

Broader Transfer Learning Success: - Cross-domain prediction often works better than expected - Suggests universal patterns in software defect manifestation - Questions assumptions about domain-specific modeling requirements

Menzies's 4th Law:

"For SE, the best thing to do with most data is to throw it away."

Supporting Evidence: - **Chen, Kocaguneli, Tu, Peters, and Xu et al.** findings across multiple prediction tasks: - **Github issue close time:** Ignored 80% of data labels [Che19] -

Effort estimation: Ignored 91% of data [Koc13] - **Defect prediction:** Ignored 97% of data [Pet15] - **Some tasks:** Ignored 98-100% of data [Che05] - **Startling result:** Data sets with thousands of rows modeled with just **few dozen samples** [Men08]

Theoretical Explanations: - **Power laws** in software data [Lin15] - **Large repeated structures** in SE projects [Hin12] - **Manifold assumption** and **Johnson-Lindenstrauss lemma** [Zhu05], [Joh84]

Caveat: Applies to regression, classification, optimization - generative tasks may still need massive data

Menzies's 5th Law:

"Bigger is not necessarily better."

Supporting Evidence - LLM Hype Analysis: - **Systematic review [Hou24]:** 229 SE papers using Large Language Models - **Critical finding:** Only 13/229 around 5% compared LLMs to other approaches - *"Methodological error"* - other PROMISE-style methods often better/faster [Gri22], [Som24], [Taw23], [Maj18]

Tree-based Models Superiority: - Grinsztajn et al. [Gri22]: *"Why do tree-based models still outperform deep learning on typical tabular data?"* - Johnson & Menzies [Joh24]: *"AI over-hype: A dangerous threat (and how to fix it)"*

Trading Off Complexity: - Scalability vs. privacy vs. performance [Lin24], [Fu17] - Often simpler methods provide better cost-effectiveness - **Personal Pattern:** *"Year later, I have switched to the simpler approach"* [Agr21], [Tan16], [Fu16]

Menzies's 6th Law:

"Data quality matters less than you think."

Supporting Research: - Shepperd et al. [She13]: Found numerous PROMISE data quality issues - Repeated rows, illegal attributes, inconsistent formats - **Critical gap:** Never tested if quality issues decreased predictive power

Our Experiment: - Built **mutators** that injected increasing amounts of their quality issues into PROMISE defect datasets - **Startling result:** Performance curves remained **flat** despite increased quality problems - **Implication:** *"There is such a thing as too much care"* in data collection

Practical Impact: - Effective predictions possible from seemingly dirty data - Questions excessive data cleaning efforts in SE research - Balance needed: careful collection without over-engineering

Menzies's 7th Law:

"Bad learners can make good conclusions."

Supporting Evidence: - Nair et al. [Nai17]: CART trees built for multi-objective optimization - **Key finding:** Models that **predicted poorly** could still **rank solutions effectively** - Could be used to prune poor configurations and find better ones - **Implication:** Algorithms shouldn't aim for predictions but offer **weak hints** about project data

Menzies's 8th Law:

"Science has mud on the lens."

Supporting Evidence: - **Hyperparameter optimization** lessons [Agr21], [Tan16], [Fu16] on PROMISE data - Data mining conclusions **changeable in an afternoon** by grad student with sufficient CPU - **Critical Questions:** Are all conclusions brittle? How build scientific community on such basis? - **Where are stable conclusions** for building tomorrow's ideas? **Bayesian Approach Needed:** Address uncertainty quantification and robust foundations

Menzies's 9th Law:

"Many hard SE problems, aren't."

Supporting Philosophy: - **Cohen's Straw Man Principle [Coh95]:** *"Supposedly sophisticated methods should be benchmarked against seemingly stupider ones"*

Personal Experience Pattern: - *"Whenever I checked a supposedly sophisticated method against a simpler one, there was always something useful in the simpler" - "More often than not, a year later, I have switched to the simpler approach"* [Agr21], [Tan16], [Fu16]

Important Caveat: - **Not all SE problems can be simplified** - Generation tools probably need LLM complexities - Safety-critical software certification remains complex - *"Just because some tasks are hard, does not mean all tasks are hard"*

Challenge to Community: *"Have we really checked what is really complex and what is really very simple?"*

Current Focus: Minimal data approaches - landscape analysis [Che19], [Lus24], surrogate learning [Nai20], active learning [Kra15], [Yu18]

PROMISE Revival Strategy (Gema Rodríguez-Pérez): - Data sharing now expected for almost all SE papers - PROMISE must differentiate: accept higher quality datasets - Focus on enhancing current data space, conducting quality evaluations

Steffen Herbold's Caution: - Early PROMISE: Collections of metrics (not raw data) - MSR shift: Raw data + fast tools (e.g., PyDriller, GHTorrent) - **Risk:** *"Little curation, little validation, often purely heuristic data collection without quality checks"* [Her22]

Modern Data Access: 1100+ recent Github projects [Xia22], CommitGuru [Ros15]

Contemporary Approaches: - **DeepLineDP** (Pornprasit et al. [Por23]): Deep learning for line-level defect prediction (TSE Best Paper 2023) - **Model interpretability:** Growing research focus [Tan21] - **Multi-objective optimization:** Hyperparameter selection [Xia22], unfairness reduction [Cha20], [Alv23]

CPU-Intensive Algorithms: - MaxWalkSat [Men09] - Simulated annealing [Men02], [Men07]

- Genetic algorithms

Minimal Data Approaches: - How much can be achieved with as little data as possible? - Suspicion of “large number of good quality labels” assumption

Cross-Domain Success [Tur09]: - **Turkish white goods** → **NASA systems** error prediction - Expected: Complex multi-dimensional transforms - **Reality:** Simple nearest neighboring between test and training data

Implication: *“Many distinctions made about software are spurious and need to be revisited”*

Power Laws & Repeated Structures: - **Lin & Whitehead [Lin15]:** Fine-grained code changes follow power laws - **Hindle et al. [Hin12]:** Software naturalness - large repeated structures - **Result:** Thousands of rows modeled with few dozen samples **[Men08]**

Lessons Learned: 1. **Open science communities** can be formed by publishing baseline + data + scripts 2. **Reproducible research** drives field advancement when embraced collectively 3. **Simple solutions** often outperform sophisticated ones 4. **Data quality** matters less than expected for predictive tasks 5. **Transfer learning** works across surprisingly diverse domains

Call-to-Action: - *“Have we really checked what is really complex and what is really very simple?”* - Challenge assumptions about problem complexity - Benchmark sophisticated methods against simpler alternatives - Focus on stable, reproducible conclusions

Slide 20: Key References (Part 1: A-K)

[Agr18]: A. Agrawal and T. Menzies, “Is better data better than better data miners?: On the benefits of tuning smote for defect prediction,” in *Proc. IST*, ACM, 2018, pp. 1050–1061. **[Agr21]:** A. Agrawal *et al.*, “How to “DODGE” complex software analytics?” *IEEE Trans. Softw. Eng.*, vol. 47, no. 10, pp. 2182–2194, Oct. 2021. **[Alv23]:** L. Alvarez and T. Menzies, “Don’t lie to me: Avoiding malicious explanations with STEALTH,” *IEEE Softw.*, vol. 40, no. 3, pp. 43–53, May/Jun. 2023. **[Cha20]:** J. Chakraborty *et al.*, “Fairway: A way to build fair ML software,” in *Proc. FSE*, 2020, pp. 654–665. **[Che19]:** J. Chen *et al.*, “‘Sampling’ as a baseline optimizer for search-based software engineering,” *IEEE Trans. Softw. Eng.*, vol. 45, no. 6, pp. 597–614, Jun. 2019. **[Coh95]:** P. R. Cohen, *Empirical Methods for Artificial Intelligence*, Cambridge, MA: MIT Press, 1995. **[Dev15]:** P. Devanbu, “Foreword,” in *Sharing Data and Models in Software Engineering*, T. Menzies *et al.*, Eds. San Mateo, CA: Morgan Kaufmann, 2015, pp. vii–viii. **[Fu16]:** W. Fu, T. Menzies, and X. Shen, “Tuning for software analytics: Is it really necessary?” *Inform. Softw. Technol.*, vol. 76, pp. 135–146, 2016. **[Gon23]:** J. M. Gonzalez-Barahona and G. Robles, “Revisiting the reproducibility of empirical software engineering studies based on data retrieved from development repositories,” *Inf. Softw. Technol.*, vol. 164, 2023, Art. no. 107318. **[Gri22]:** L. Grinsztajn, E. Oyallon, and G. Varoquaux, “Why do tree-based models still outperform deep learning on typical tabular data?” in *Proc. NeurIPS*, 2022, pp. 507–520. **[Ham09]:** M. Hamill and K. Goseva-Popstojanova, “Common trends in software fault and failure data,” *IEEE Trans. Softw. Eng.*, vol. 35, no. 4, pp. 484–496, Jul./Aug. 2009. **[Has08]:** A. E. Hassan, “The road ahead for mining software repositories,” *Frontiers Softw. Maintenance*, pp. 48–57, 2008.

Slide 21: Key References (Part 2: K-Z)

[Kim15]: M. Kim *et al.*, “REMI: Defect prediction for efficient api testing,” in *Proc. FSE*, ACM, 2015, pp. 990–993. **[Kri19]:** R. Krishna and T. Menzies, “Bellwethers: A baseline method for transfer learning,” *IEEE Trans. Softw. Eng.*, vol. 45, no. 11, pp. 1081–1105, Nov. 2019. **[Men07]:** T. Menzies, J. Greenwald, and A. Frank, “Data mining static code attributes to learn defect predictors,” *IEEE Trans. Softw. Eng.*, vol. 33, no. 1, pp. 2–13, Jan. 2007. **[Men24]:** T. Menzies, “A brief note, with thanks, on the contributions of guenther ruhe,” *Inf. Softw. Technol.*, vol. 173, 2024, Art. no. 107486. **[Mis11]:** A. T. Misirli, A. Bener, and R. Kale, “AI-based software defect predictors: Applications and benefits in a case study,” *AI Mag.*, vol. 32, no. 2, pp. 57–68, 2011. **[Nai17]:** V. Nair *et al.*, “Using bad learners to find good configurations,” in *Proc. 11th Joint Meeting FSE*, ACM, 2017, pp. 257–267. **[Nam18]:** J. Nam *et al.*, “Heterogeneous defect prediction,” *IEEE Trans. Softw. Eng.*, vol. 44, no. 9, pp. 874–896, Sep. 2018. **[Ost04]:** T. J. Ostrand, E. J. Weyuker, and R. M. Bell, “Where the bugs are,” *ACM SIGSOFT Softw. Eng. Notes*, vol. 29, no. 4, pp. 86–96, 2004. **[Por23]:** C. Pornprasit and C. K. Tantithamthavorn, “DeepLineDP: Towards a deep learning approach for line-level defect prediction,” *IEEE Trans. Softw. Eng.*, vol. 49, no. 1, pp. 84–98, Jan. 2023. **[Rah14]:** F. Rahman *et al.*, “Comparing static bug finders and statistical prediction,” in *Proc. ICSE*, ACM, 2014, pp. 424–434. **[Rob10]:** G. Robles, “Replicating MSR: A study of the potential replicability of papers published in the mining software repositories proceedings,” in *7th IEEE Work. Conf. Mining Softw. Repositories (MSR)*, IEEE Press, 2010, pp. 171–180. **[Ros15]:** C. Rosen, B. Grawi, and E. Shihab, “Commit guru: Analytics and risk prediction of software commits,” in *Proc. ESEC/FSE 2015*, pp. 966–969. **[She13]:** M. Shepperd *et al.*, “Data